ANGEL B

AI Engineer | LLM Applications & Document Intelligence

Bengaluru, India | Open to relocation to UAE | +91 9567495787 | asha4angelsb@gmail.com

LinkedIn | GitHub

Availability: 30 Nov 2026 | Visa: Requires employment visa sponsorship

PROFESSIONAL SUMMARY

AI Engineer with 1 year of production experience building LLM-powered document intelligence systems. Lead developer in a 3-person team on an end-to-end platform that classifies industrial specification documents, extracts structured requirements with LLMs, and generates print-ready outputs with human-in-the-loop review. Experienced in multi-provider LLM integration (Google Gemini, OpenAI), PostgreSQL-backed pipelines, and containerized deployment, with a classical ML foundation (XGBoost, LightGBM, CatBoost). Seeking AI/GenAI engineering roles in the UAE.

TECHNICAL SKILLS

Languages: Python, SQL

LLM & GenAI: Google Gemini API, OpenAI GPT API, prompt engineering, multimodal extraction, structured output, LLM workflow automation, human-in-the-loop review design

Document AI & Vision: OCR (PaddleOCR, Tesseract), OpenCV, PyZbar, document classification and structured extraction

Machine Learning: Scikit-learn, XGBoost, LightGBM, CatBoost, HDBSCAN clustering, feature engineering, SMOTE

Data & Backend: PostgreSQL, SQLAlchemy, database schema design, pipeline design, REST API integration

Tools & Deployment: Docker, Git/GitHub, Pytest, Streamlit

Currently Learning: RAG, embeddings and vector databases, FastAPI, LangGraph, cloud deployment

PROFESSIONAL EXPERIENCE

AI Engineer | Digitus Business Solutions Pvt. Ltd., Bengaluru *Nov 2025 – Present*

- Lead developer in a 3-person team on an industrial document intelligence and label-generation platform, owning work from architecture and database design through backend services, review tooling, and Docker deployment.
- Designed a document classification module that sorts uploaded specifications into industrial, global, and client-specific standards before storing them in PostgreSQL, making the platform reusable across industries instead of a single client.
- Built an LLM extraction pipeline that converts industrial documents and guidelines into structured tables of label variants, requirements, and reference images, with a human approval step before data is saved.
- Built a label generation engine that produces previews from extracted data and reference images, with two editable outputs (visual canvas and raw ZPL code) for fine-tuning before export.
- Designed the initial image-only input flow and conversational interface for creating labels from a single reference image; the flow was later extended by the team with an OCR module.
- Implemented dimension- and DPI-aware rendering and verified physical print output matches the specified size and DPI.
- Reduced label design and confirmation turnaround by over 95% (from 2 days to 10-20 minutes) by automating extraction, standards matching, and generation.
- Built a template reuse workflow that stores approved layouts, cutting repeat LLM API calls, cost, and turnaround for minor client change requests.
- Leading migration of the review and canvas-editing interface from Streamlit to a full web application, and leading handover to the incoming engineering team through architecture, workflow, and codebase documentation.
- Led and coordinated two engineers who built the OCR module and added canvas editor tools, while owning architecture, the LLM pipeline, and backend services.
- Built smaller internal ML and document-processing tools for ad hoc automation needs.

Data Science & AI Intern | Rubixe, Bengaluru *Feb 2025 – Aug 2025*

- Built and evaluated ML pipelines (XGBoost, LightGBM, CatBoost) for classification and regression on clinical, financial, and automotive datasets, including IQR outlier treatment, feature engineering, and SMOTE class balancing.
- Delivered an import car price estimator ($R^2 = 0.94$) and a customer transaction prediction model (ROC-AUC 0.81).

KEY PROJECT

Industrial Document-to-ZPL AI Label Generation Platform | *Digitus Business Solutions (Ongoing)*

Tech: Python, Google Gemini API, OpenAI GPT, PaddleOCR, OpenCV, PostgreSQL, SQLAlchemy, Streamlit, Docker

- Originated the project by generalizing a single-client automation request into a configurable solution for multiple industries and client standards.

- Designed and owned the architecture covering classification, schema, LLM extraction, label generation, canvas/ZPL export, and print-ready rendering.
- Built the human-in-the-loop review interface so users can edit extracted fields and generated labels before approved data is saved to PostgreSQL.
- Processed 20+ industrial specification documents and generated 100+ labels for manufacturing and automotive use cases, with printed dimensions verified against a physical scale.
- Achieved 90%+ extraction accuracy on specification documents; storing approved extractions and labels removes re-extraction on repeat requests, reducing LLM API cost.

EDUCATION

B.Tech, Computer Science and Engineering *2019 – 2023*
Government College of Engineering, Kannur

CERTIFICATIONS

- Data Science and Artificial Intelligence Expert, IABAC (Jul 2025)
- Certified Data Scientist, NASSCOM (Nov 2024)
- Certified Specialist in Machine Learning AI, ICT Academy of Kerala (Feb 2024)

LANGUAGES

English (fluent), Malayalam (native), Hindi (intermediate), Kannada and Tamil (basic understanding)